

Information Geometry: The Fisher Metric

math-foundations / concept 40

1 Setting

Let $\{p_\theta : \theta \in \Theta\}$ be a smoothly parameterized family of probability densities on a measurable space $\mathcal{X}$, with $\Theta \subseteq \mathbb{R}^d$ open. We assume sufficient regularity to interchange differentiation in θ with integration in x , and that $\log p_\theta(x)$ is twice continuously differentiable in θ for almost every x . Define the *score* $s_\theta(x) := \nabla_\theta \log p_\theta(x)$.

Definition 1 (Fisher information matrix). *The Fisher information matrix at θ is the $d \times d$ matrix*

$$F(\theta) = \mathbb{E}_{X \sim p_\theta} [s_\theta(X) s_\theta(X)^\top] = -\mathbb{E}_{X \sim p_\theta} [\nabla_\theta^2 \log p_\theta(X)].$$

The two expressions agree under the regularity assumptions above.

Definition 2 (Statistical manifold). *A statistical manifold is a pair (Θ, F) where Θ is a smooth manifold parameterizing $\{p_\theta\}$ and F is the Fisher information matrix viewed as a Riemannian metric tensor: for $u, v \in T_\theta \Theta$, $\langle u, v \rangle_\theta := u^\top F(\theta) v$.*

Definition 3 (Natural gradient). *Given a smooth loss $L : \Theta \rightarrow \mathbb{R}$, the natural gradient is*

$$\tilde{\nabla} L(\theta) = F(\theta)^{-1} \nabla L(\theta),$$

defined wherever $F(\theta)$ is positive definite. It is the steepest-descent direction under the Fisher metric.

2 The score has mean zero

Lemma 1. $\mathbb{E}_{X \sim p_\theta} [s_\theta(X)] = 0$.

Proof. Since $\int p_\theta(x) dx = 1$ for every θ , differentiating in θ and exchanging derivative and integral,

$$0 = \nabla_\theta \int p_\theta(x) dx = \int \nabla_\theta p_\theta(x) dx = \int \frac{\nabla_\theta p_\theta(x)}{p_\theta(x)} p_\theta(x) dx = \mathbb{E}_{p_\theta} [s_\theta(X)]. \quad \square$$

3 Fisher is the local KL metric

Theorem 1. *The second-order Taylor expansion of $\epsilon \mapsto D_{\text{KL}}(p_\theta \| p_{\theta+\epsilon})$ near $\epsilon = 0$ is*

$$D_{\text{KL}}(p_\theta \| p_{\theta+\epsilon}) = \frac{1}{2} \epsilon^\top F(\theta) \epsilon + O(\|\epsilon\|^3).$$

Proof. Define $\Phi(\epsilon) := D_{\text{KL}}(p_\theta \| p_{\theta+\epsilon}) = \mathbb{E}_{X \sim p_\theta} [\log p_\theta(X) - \log p_{\theta+\epsilon}(X)]$. Set $\ell(\epsilon, x) := \log p_{\theta+\epsilon}(x)$, so $\Phi(\epsilon) = \mathbb{E}_{p_\theta} [\ell(0, X) - \ell(\epsilon, X)]$.

Taylor-expand $\ell(\epsilon, x)$ in ϵ at 0 to second order, with remainder $R(\epsilon, x) = O(\|\epsilon\|^3)$ uniformly on compact sets under our smoothness hypotheses:

$$\ell(\epsilon, x) = \ell(0, x) + \epsilon^\top \nabla_\epsilon \ell(0, x) + \frac{1}{2} \epsilon^\top \nabla_\epsilon^2 \ell(0, x) \epsilon + R(\epsilon, x).$$

Note $\nabla_\epsilon \ell(0, x) = s_\theta(x)$ and $\nabla_\epsilon^2 \ell(0, x) = \nabla_\theta^2 \log p_\theta(x)$. Substituting,

$$\ell(0, x) - \ell(\epsilon, x) = -\epsilon^\top s_\theta(x) - \frac{1}{2} \epsilon^\top \nabla_\theta^2 \log p_\theta(x) \epsilon - R(\epsilon, x).$$

Take expectation under p_θ . By Lemma 1, $\mathbb{E}_{p_\theta} [\epsilon^\top s_\theta(X)] = \epsilon^\top \mathbb{E}_{p_\theta} [s_\theta(X)] = 0$, so the linear term vanishes. By the Hessian form of Fisher, $-\mathbb{E}_{p_\theta} [\nabla_\theta^2 \log p_\theta(X)] = F(\theta)$, so

$$\Phi(\epsilon) = 0 + \frac{1}{2} \epsilon^\top F(\theta) \epsilon - \mathbb{E}_{p_\theta} [R(\epsilon, X)] = \frac{1}{2} \epsilon^\top F(\theta) \epsilon + O(\|\epsilon\|^3). \quad \square$$

Consequence. The Hessian of $\epsilon \mapsto D_{\text{KL}}(p_\theta \| p_{\theta+\epsilon})$ at $\epsilon = 0$ is $F(\theta)$. The Fisher information is therefore the unique tensor that captures the local quadratic behaviour of KL divergence — justifying its role as the natural Riemannian metric on the statistical manifold and as the preconditioner in natural-gradient descent.